Lightweight Deferred Updates for Linux Kenrel Scalability

```
function logical_insert(obj, root):
  If CAS(obj.del_node.mark, 1, 0) ≠ 1:
    obj.add_node.mark ← 1
    If test_and_set_bit(OP_INSERT, obj.exist) ≠ true:
      set_bit(OP_INSERT, obj.used):
      obj.add_node.op ← OP_INSERT
      obj.add_node.key ← obj
      obj.add_node.root ← root
      add_lock_less_list(obj.add_node)

function logical_remove(obj, root):
  If CAS(obj.add_node.mark, 1, 0) ≠ 1:
    obj.del_node.mark ← 1
    If test_and_set_bit(OP_REMOVE, obj.exist) ≠ true:
      set_bit(OP_REMOVE, obj.used):
      obj.del_node.op ← OP_REMOVE
      obj.del_node.key ← obj
      obj.del_node.root ← root
      add_lock_less_list(obj.del_node)
```

Figure 1: LDU logical update algorithm. `logical_insert` represents non-blocking insert function. It may be called by original insert position without locks. The fastpath is that when their object was removed by `logical_remove`, `logical_insert` just changes node's marking field.

```
function synchronize_ldu(obj, head):
  If (head.first = NULL):
    return;
  entry ← xchg(head.first, NULL);
  for each list node:
    obj ← node.key
    clear_bit(node.op, obj.exist)
    If !xchg(node.mark, 0):
        physical_update(node.op, obj, node.root)
    clear_bit(node.op, obj.used)

function physical_update(op, obj, root):
  If op = OP_INSERT :
    call real insert function(obj, root)
  Else If op = OP_REMOVE :
    call real remove function(obj, root)
```

Figure 2: LDU physical update algorithm. `synchronize_ldu` may be called by reader and converts update log to original data structure traversing the lock-less list.


```
function logical_insert(obj, root):
  If CAS(obj.del_node.mark, 1, 0) ≠ 1:
    obj.add_node.mark ← 1
    If test_and_set_bit(OP_INSERT, obj.exist) ≠ true:
      set_bit(OP_INSERT, obj.used):
      obj.add_node.op ← OP_INSERT
      obj.add_node.key ← obj
      obj.add_node.root ← root
      add_lock_less_list(obj.add_node)

function logical_remove(obj, root):
  If CAS(obj.add_node.mark, 1, 0) ≠ 1:
    obj.del_node.mark ← 1
    If test_and_set_bit(OP_REMOVE, obj.exist) ≠ true:
      set_bit(OP_REMOVE, obj.used):
      obj.del_node.op ← OP_REMOVE
      obj.del_node.key ← obj
      obj.del_node.root ← root
      add_lock_less_list(obj.del_node)
```

Figure 3: LDU logical update algorithm. `logical_insert` represents non-blocking insert function. It may be called by original insert position without locks. The fastpath is that when their object was removed by `logical_remove`, `logical_insert` just changes node's marking field.

```
function synchronize_ldu(obj, head):
  If (head.first = NULL):
    return;
  entry ← xchg(head.first, NULL);
  for each list node:
    obj ← node.key
    clear_bit(node.op, obj.exist)
    If !xchg(node.mark, 0):
        physical_update(node.op, obj, node.root)
    clear_bit(node.op, obj.used)

function physical_update(op, obj, root):
  If op = OP_INSERT :
    call real insert function(obj, root)
  Else If op = OP_REMOVE :
    call real remove function(obj, root)
```

Figure 4: LDU physical update algorithm. `synchronize_ldu` may be called by reader and converts update log to original data structure traversing the lock-less list.